

CSE 5693 Machine Learning HW2

Due 6:30pm, Feb 24

Submit Server: Class = cse5693 , Assignment = hw2

1. Written assignment:

- (a) 2.4
- (b) 2.7
- (c) 3.4
- (d) Consider two attributes Outlook (sunny, rainy, cloudy) and Humidity (high) and outcome PlayTennis (yes, no) for the instance space (X).
 - i. Consider an unbiased hypothesis space ($H1$), enumerate all possible “Yes” hypotheses (h_1, h_2, \dots) in terms of subsets of instances. What is the number of possible unique hypotheses in $H1$?
 - ii. For each hypothesis in $H1$, represent it as a boolean expression. What is the number of unique hypotheses semantically?
 - iii. Consider a biased hypothesis space ($H2$) where each attribute can only have a value, ?, or \emptyset . What is the number of unique hypotheses in the biased hypothesis space ($H2$)?
 - iv. Identify hypotheses in the unbiased hypothesis space ($H1$) that are not in the biased hypothesis space ($H2$).
- (e) With the programming assignment: Discuss and compare accuracy of no pruning versus rule post-pruning in testTennis, testIris, and testIrisNoisy.

2. Programming assignment: Decision Tree

- (a) Allow more than two outcomes/classes
- (b) Allow continuous-valued attributes
- (c) Allow printing the tree
- (d) Allow the option of rule post-pruning and printing the rules
- (e) Two data sets: Tennis and Iris on the course web site.
- (f) The same program should be able to handle the two data sets.
- (g) For each of the following experiments, provide a script/program/function to run the experiment:
 - i. testTennis: print the tree, accuracy on the training and test sets, the rules after post-pruning, accuracy on the training and test sets
 - ii. testIris: print the tree, accuracy on the training and test sets, the rules after post-pruning, accuracy on the training and test sets

- iii. testIrisNoisy: corrupt the class labels of training examples from 0% to 20% (2% increment) by changing from the correct class to another class; output the accuracy on the uncorrupted test set with and without rule post-pruning.

(h) Implementation:

- i. Use C (GNU gcc), C++ (GNU g++), Java (Oracle Java), LISP (CLISP), or Python. If you don't have a preference, use Java since it's more portable.
- ii. Your program preferably runs on code01.fit.edu (linux).
- iii. You might have these modules:

- A. Learner: input training examples/instances, output a tree (or rule set)
- B. Classifier/predictor: input a tree (or rule set) and labeled instances, output the classifications/predictions and how accurate the tree is with respect to the correct labels (% of correct classifications).
- C. Tree printer (pre-order traversal, deeper nodes are indented more, leaves have class distribution), for example:

```
height = tall
|   size>2 = T
|   |   color = black
|   |   |   weight = heavy : Yes (1,0)
|   |   |   weight = light : No (0,1)
|   |   color = white
|   |   |   weight = heavy : Yes (2,0)
|   |   |   weight = light : No (0,1)
|   size>2 = F
|   |   weight = heavy : Yes (4,0)
|   |   weight = light : No (0,2)
height = short : No (0,8)
```

- D. Rule set printer, for example:

```
height = tall ^ size>2 = T => Yes (1,0)
height = tall ^ size>2 = F => No (0,1)
```

iv. Submission:

- A. README.txt: what are the files and how to compile and run your program on code01.fit.edu
- B. source code